

Manipulating Text Files

- Irrespective of the role you play with Linux (system administrator, developer or user), you often need to browse through and parse text files, and/or extract data from them
- These are **file manipulation operations** and it is essential for the Linux user to become adept at performing certain operations on files
- Most of the time, such file manipulation is done at the command line, which allows users to perform tasks more efficiently than while using a GUI; the command line is also more suitable for automating often executed tasks
- Experienced system administrators write customized scripts to accomplish such repetitive tasks, standardized for each particular environment

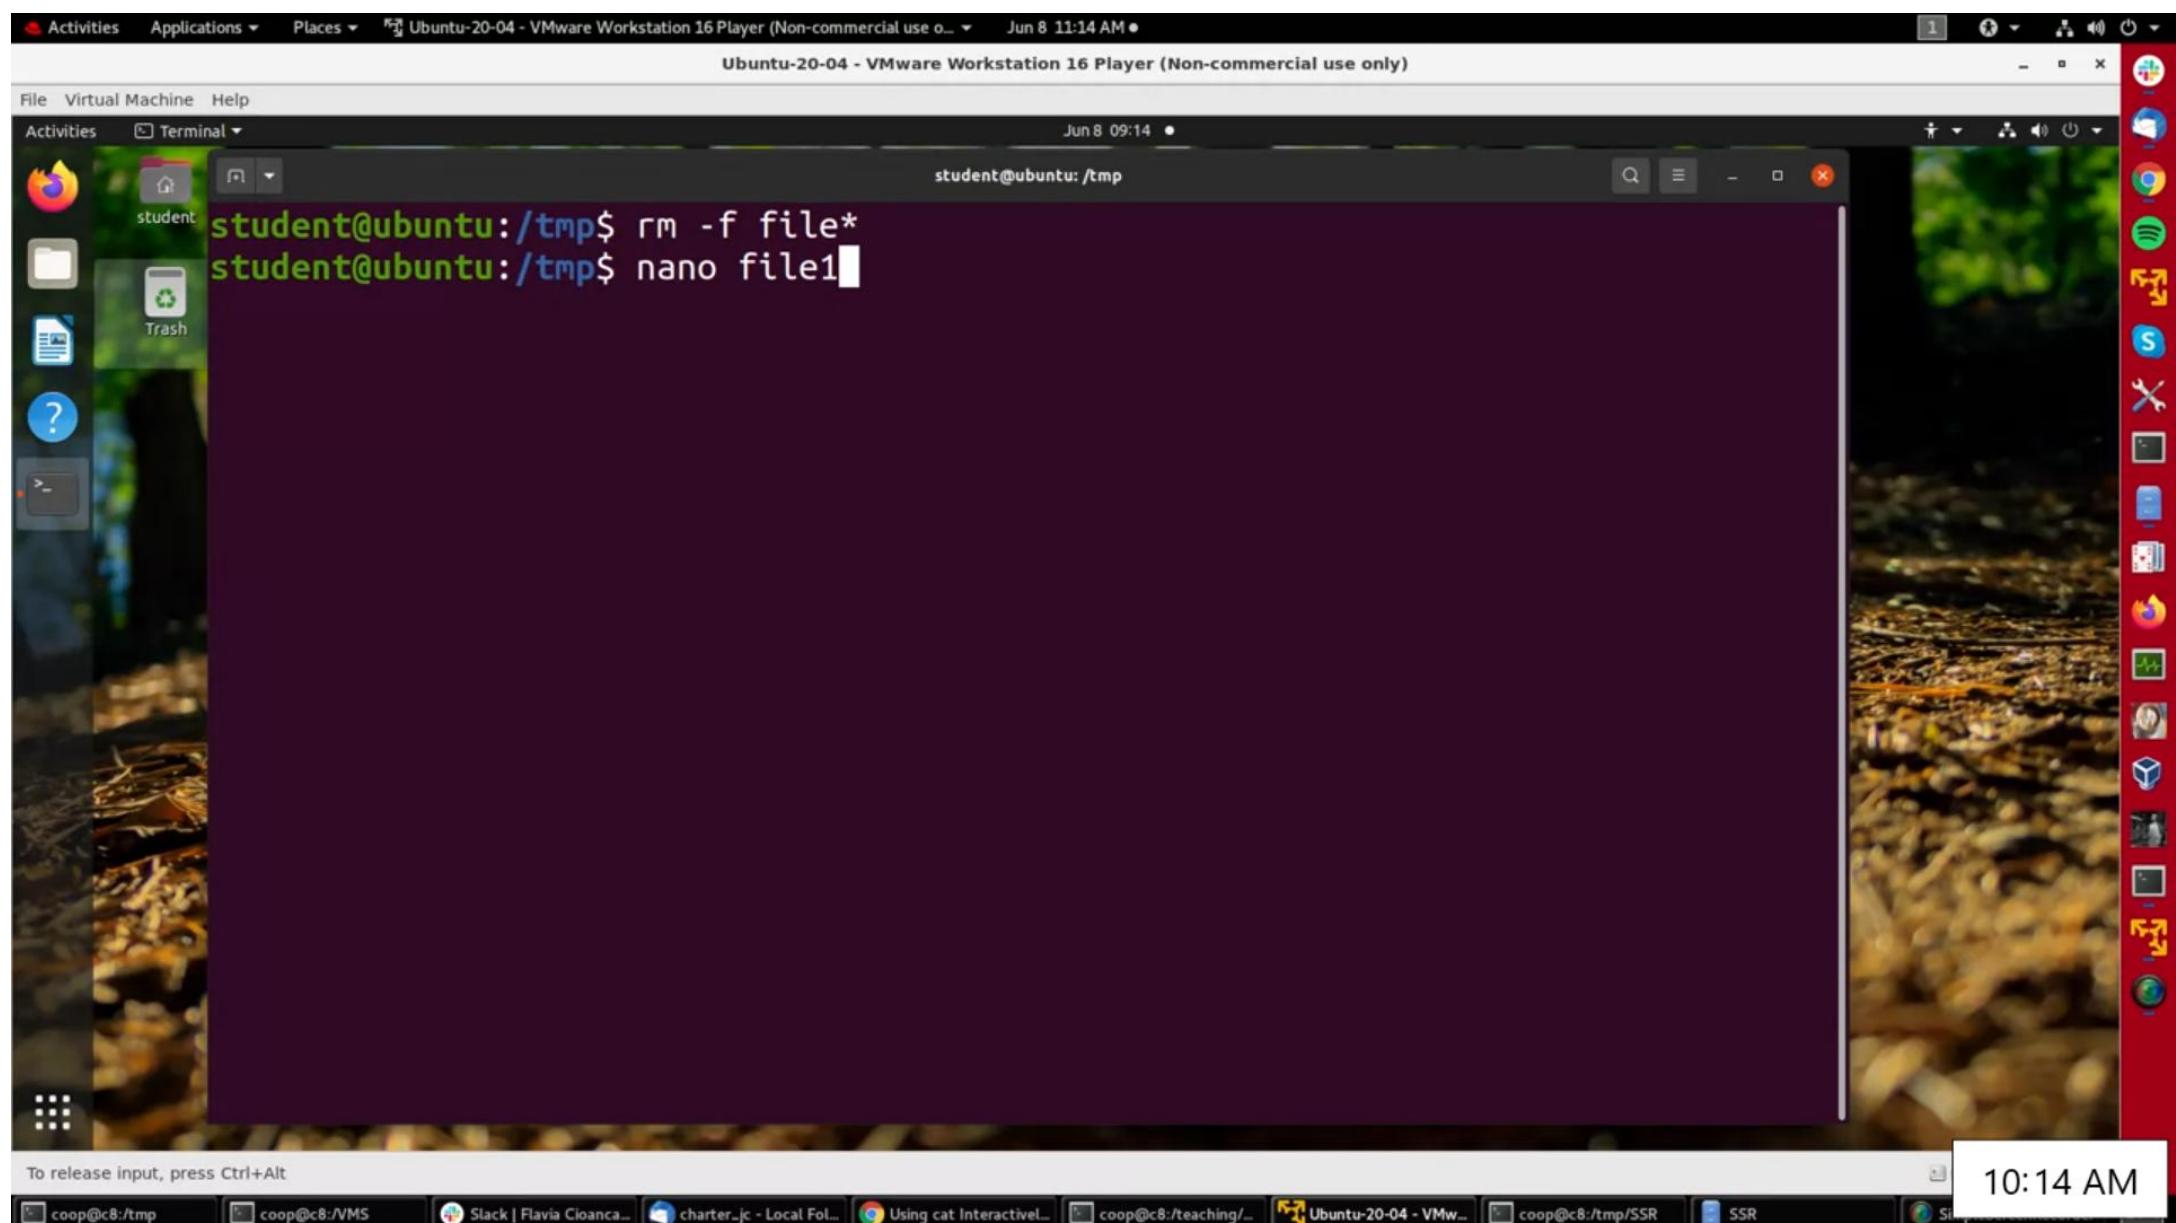
Manipulating Text Files

- Irrespective of the role you play with Linux (system administrator, developer or user), you often need to browse through and parse text files, and/or extract data from them
- These are **file manipulation operations** and it is essential for the Linux user to become adept at performing certain operations on files
- Most of the time, such file manipulation is done at the command line, which allows users to perform tasks more efficiently than while using a GUI; the command line is also more suitable for automating often executed tasks
- Experienced system administrators write customized scripts to accomplish such repetitive tasks, standardized for each particular environment

Command Line Utilities

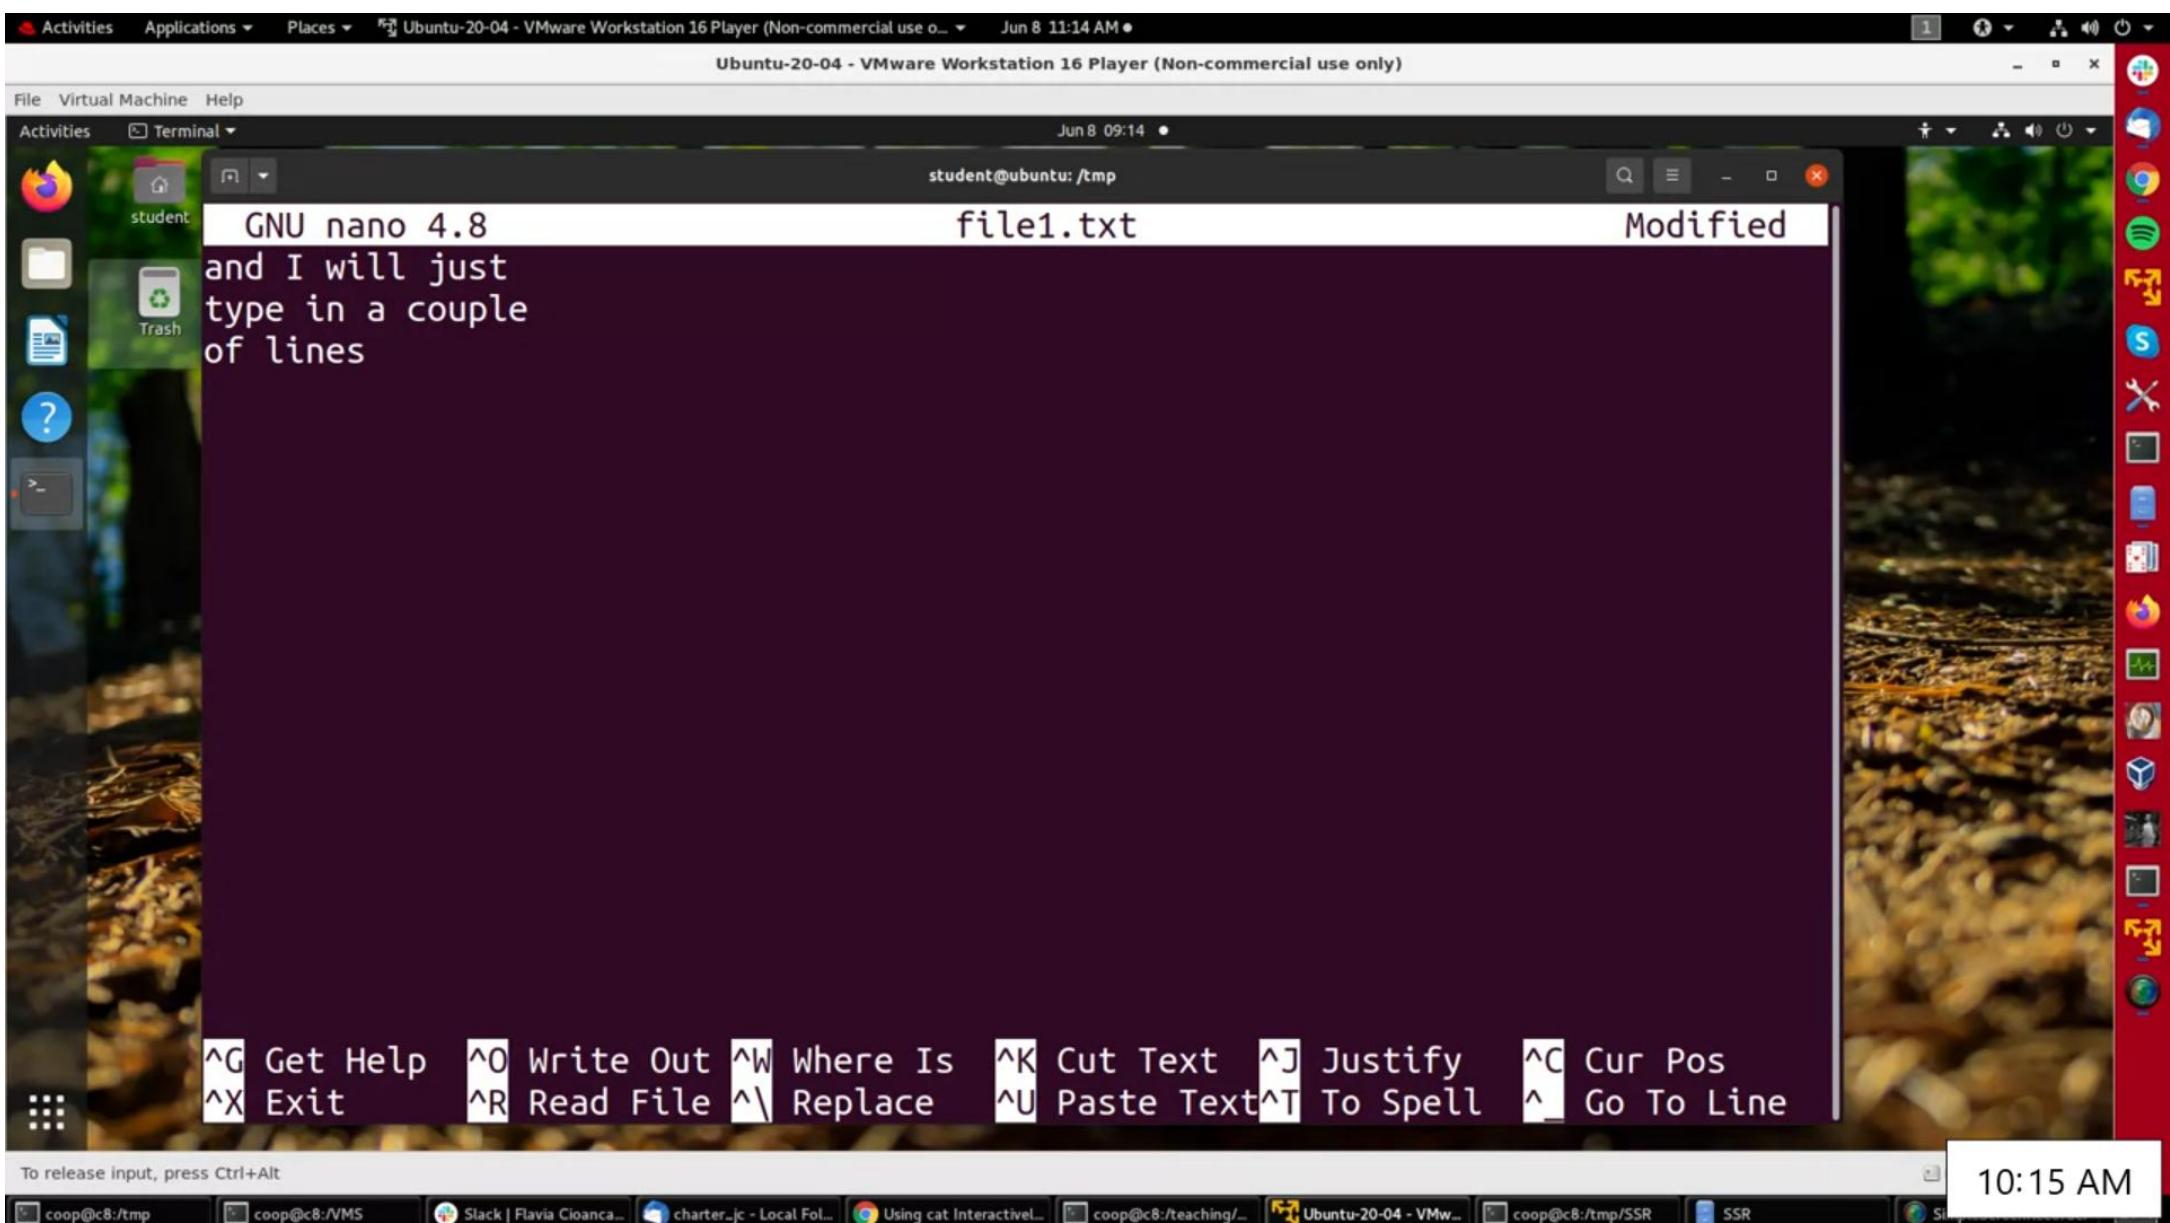
The basic Linux command line utilities are:

- **cat**
- **echo**
- **head**
- **tail**
- **sed**
- **awk**



To release input, press Ctrl+Alt

10:14 AM



Activities Applications Places Ubuntu-20-04 - VMware Workstation 16 Player (Non-commercial use o... Jun 8 11:15 AM

Ubuntu-20-04 - VMware Workstation 16 Player (Non-commercial use only)

File Virtual Machine Help

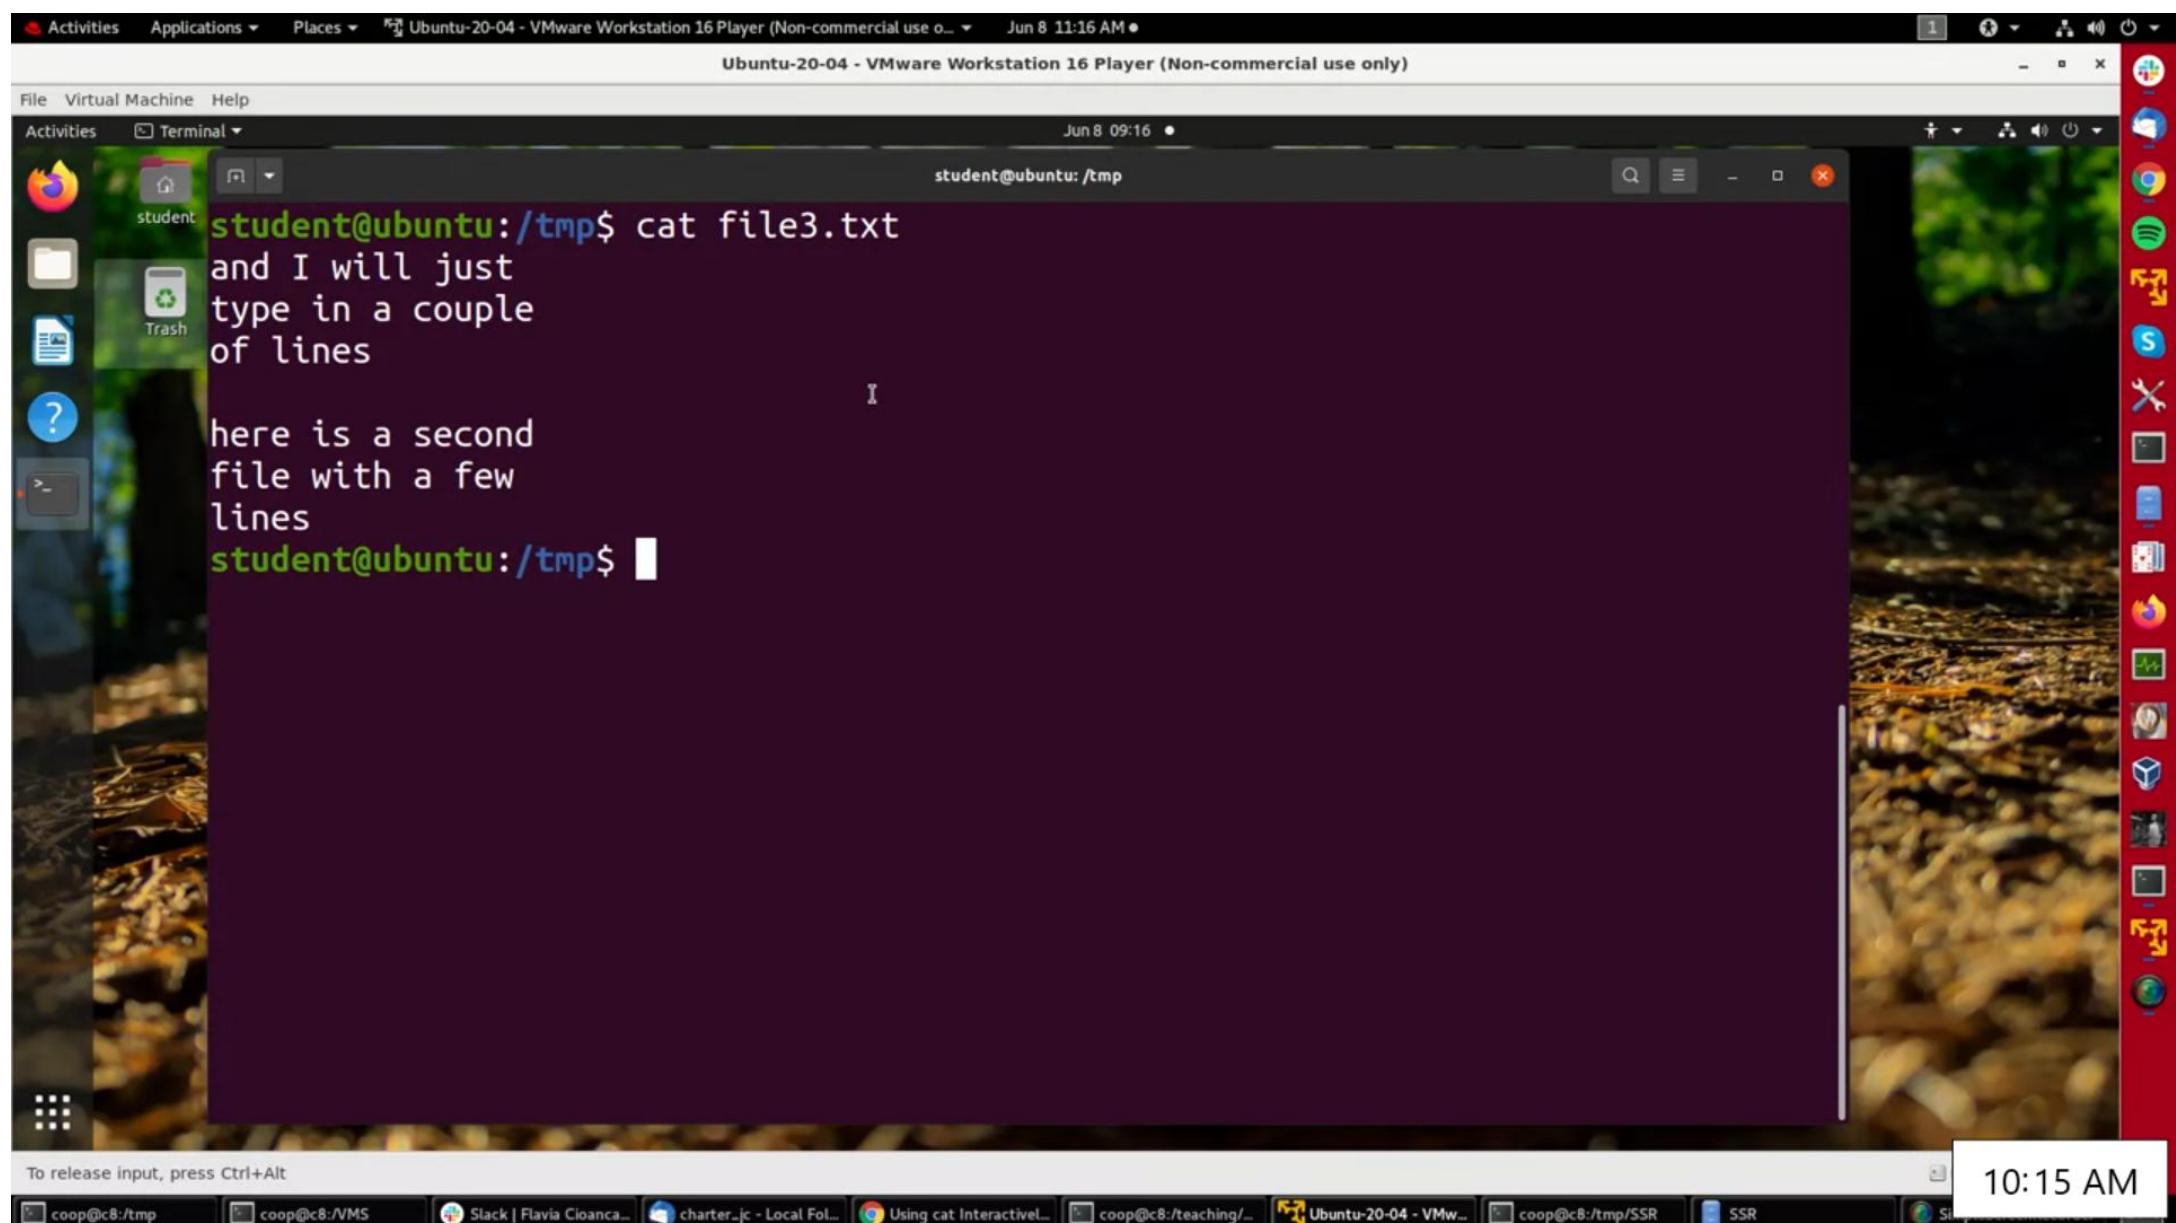
Activities Terminal Jun 8 09:15 student@ubuntu: /tmp

```
student@ubuntu:/tmp$ rm -f file*
student@ubuntu:/tmp$ nano file1.txt
student@ubuntu:/tmp$ cat file1.txt
and I will just
type in a couple
of lines

student@ubuntu:/tmp$ cat << EOF > file2.txt
> here is a second
> file with a few
> lines
> EOF
student@ubuntu:/tmp$
```

To release input, press Ctrl+Alt 10:15 AM

coop@c8:/tmp coop@c8:/VMS Slack | Flavia Cioanca... charter_jc - Local Fol... Using cat Interactivel... coop@c8:/teaching/_... Ubuntu-20-04 - VMw... coop@c8:/tmp/SSR SSR Si



Managing Large Files

- System administrators need to work with configuration files, text files, documentation files, and log files
- Some of these files may be large or become quite large as they accumulate data with time
- These files will require both viewing and administrative updating

Example

A banking system might maintain one simple large log file to record details of all of one day's ATM transactions. Due to a security attack or a malfunction, the administrator might be forced to check for some data by navigating within the file.

In such cases, directly opening the file in an editor will cause issues, due to high memory utilization, as an editor will usually try to read the whole file into memory first.

What can be a possible solution?

less and **somefile**

- One can use **less** to view the contents of such a large file, scrolling up and down page by page, without the system having to place the entire file in memory before starting; this is much faster than using a text editor
- Viewing **somefile** can be done by typing either of the two following commands:

```
$ less somefile
```

```
$ cat somefile | less
```

- By default, man pages are sent through the **less** command
- You may have encountered the older **more** utility which has the same basic function but fewer capabilities: i.e. less is more!



Press **Esc** to exit full screen

Introduction to sed and awk



Introduction to **sed** and **awk**

- It is very common to create and then repeatedly edit and/or extract contents from a file; in order to perform such operations you can use **sed** and **awk**
- Many Linux users and administrators will write scripts using comprehensive scripting languages such as **Python** and **perl**, rather than use **sed** and **awk** (and some other utilities we will discuss later)
- Using such utilities is certainly fine in most circumstances; you should always feel free to use the tools you are experienced with
- However, the utilities that are described here are much lighter; i.e. they use fewer system resources, and execute faster
- There are situations (such as during booting the system) where a lot of time would be wasted using the more complicated tools, and the system may not even be able to run them; the simpler tools will always be needed

sed

- **sed** is a powerful text processing tool and one of the oldest, earliest and most popular UNIX utilities; its name is an abbreviation for stream editor
- It is used to modify the contents of a file, usually placing the contents into a new file, and it can also filter text, as well as perform substitutions in data streams
- How it works:
 - Data from an input source/file (or stream) is taken and moved to a working space
 - The entire list of operations/modifications is applied over the data in the working space and the final contents are moved to the standard output space (or stream)



awk

- **awk** was created at Bell Labs in the 1970s and derived its name from the last names of its authors: Alfred Aho, Peter Weinberger, and Brian Kernighan
 - Used to extract and then print specific contents of a file and is often used to construct reports
 - A powerful utility and interpreted programming language
 - Used to manipulate data files, retrieving, and processing text
 - Works well with fields (containing a single piece of data, essentially a column) and records (a collection of fields, essentially a line in a file)

```
Activities Applications Terminal File Virtual Machine Help FC-34 student@student@fedora:~/tmp
[student@fedora tmp]$ cat infile.txt
This is the first line
This is the second line
This is the third line
[student@fedora tmp]$ sed -e s/is/are/ infile.txt
Thare is the first line
Thare is the second line
Thare is the third line
[student@fedora tmp]$ sed -e s/is/are/g infile.txt
Thare are the first line
Thare are the second line
Thare are the third line
[student@fedora tmp]$ sed -e 1,2s/is/are/g infile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$
```

10:19 AM

```
Activities Applications Terminal File Virtual Machine Help FC-34 student@student@fedora:~/tmp
[student@fedora tmp]$ cat infile.txt
This is the first line
This is the second line
This is the third line
[student@fedora tmp]$ sed -e s/is/are/ infile.txt
Thare is the first line
Thare is the second line
Thare is the third line
[student@fedora tmp]$ sed -e s/is/are/g infile.txt
Thare are the first line
Thare are the second line
Thare are the third line
[student@fedora tmp]$ sed -e 1,2s/is/are/g infile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$ sed -e 1,2s:is:are:g infile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$
```

10:19 AM

```
Activities Applications Terminal File Virtual Machine Help FC-34 student@student@fedora:~/tmp
[student@fedora tmp]$ cat infile.txt
This is the first line
This is the second line
This is the third line
[student@fedora tmp]$ sed -e s/is/are/ infile.txt
Thare is the first line
Thare is the second line
Thare is the third line
[student@fedora tmp]$ sed -e s/is/are/g infile.txt
Thare are the first line
Thare are the second line
Thare are the third line
[student@fedora tmp]$ sed -e 1,2s/is/are/g infile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$ sed -e 1,2s:is:are:g infile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$ sed -e 1,2s:is:are:g > outfile.txt
^C
[student@fedora tmp]$ sed -e 1,2s:is:are:g infile.txt > outfile.txt
[student@fedora tmp]$ cat
```

10:19 AM

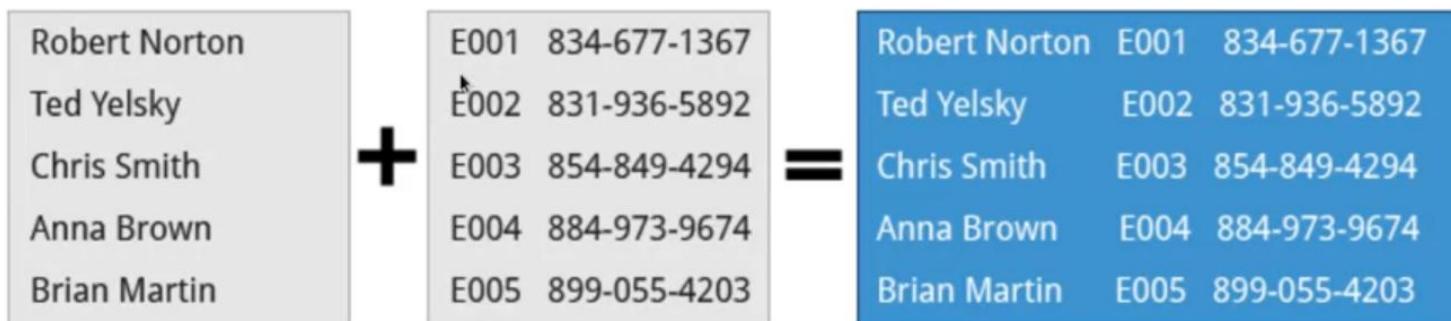
```
Activities Applications Terminal File Virtual Machine Help FC-34 student@student@fedora:tmp
[student@fedora tmp]$ cat outfile.txt
Thare are the first line
Thare are the second line
This is the third line
[student@fedora tmp]$ diff infile.txt outfile.txt
1,2c1,2
< This is the first line
< This is the second line
---
> Thare are the first line
> Thare are the second line
[student@fedora tmp]$
```

File Manipulation Utilities

- In managing your files, you may need to perform many tasks, such as sorting data and copying data from one location to another
- Linux provides several file manipulation utilities that you can use while working with text files, including:
 - **sort**
 - **uniq**
 - **paste**
 - **join**
 - **split**

Example

Suppose you have a file that contains the full name of all employees and another file that lists their phone numbers and Employee IDs. You want to create a new file that contains all the data listed in three columns: name, employee ID, and phone number. How can you do this effectively without investing too much time?



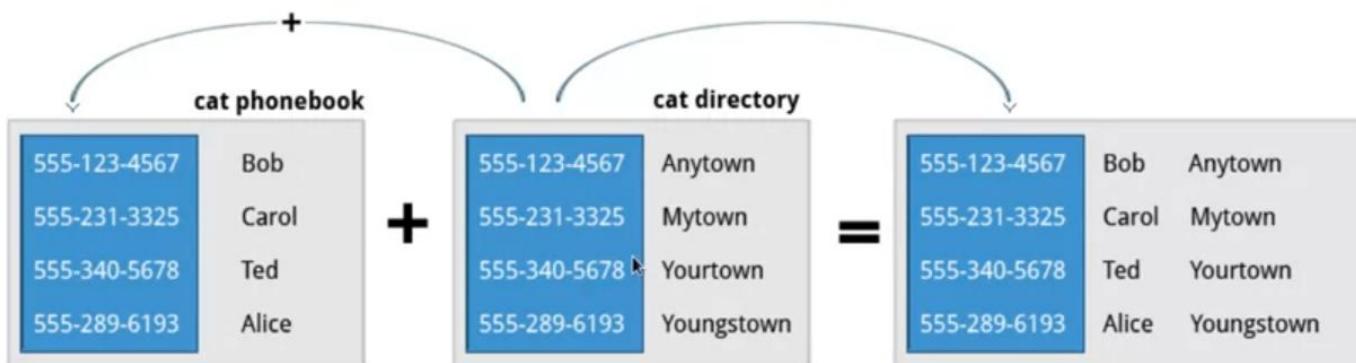
paste

- **paste** can be used to create a single file containing all three columns
- The different columns are identified based on delimiters (spacing used to separate two fields); for example, delimiters can be a blank space, a tab, or an Enter
- In the image provided, a single space is used as the delimiter in all files
- **paste** accepts the following options:
 - **-d** delimiters: specify a list of delimiters to be used instead of tabs for separating consecutive values on a single line; each delimiter is used in turn and when the list has been exhausted, **paste** begins again at the first delimiter
 - **-s** causes **paste** to append the data in series rather than in parallel; that is, in a horizontal rather than vertical fashion

Example

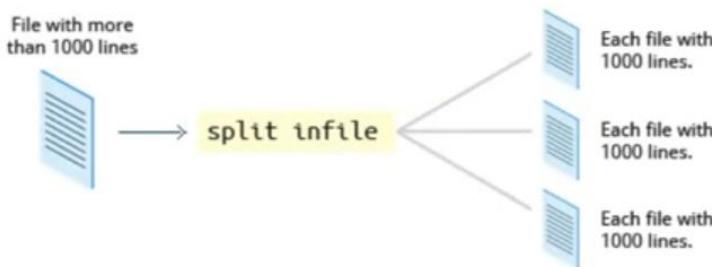
Suppose you have two files with some similar columns. You have saved employees' phone numbers in two files, one with their first name and the other with their last name. You want to combine the files without repeating the data of common columns.

How do you achieve this?



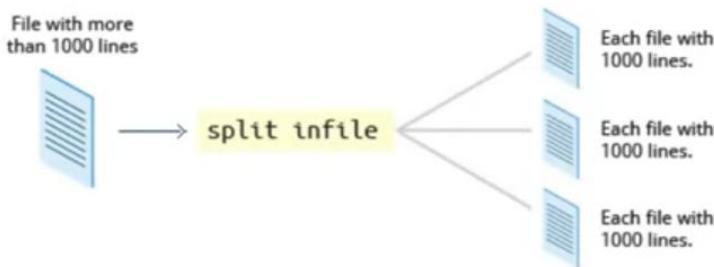
split

- **split** is used to break up (or split) a file into equal-sized segments for easier viewing and manipulation, and is generally used only on relatively large files
- By default, **split** breaks up a file into 1000-line segments; the original file remains unchanged, and a set of new files with the same name plus an added prefix is created
 - By default, the **x** prefix is added
 - To split a file into segments, use the command **split infile**
 - To split a file into segments using a different prefix, use the command **split infile <Prefix>**



split

- **split** is used to break up (or split) a file into equal-sized segments for easier viewing and manipulation, and is generally used only on relatively large files
- By default, **split** breaks up a file into 1000-line segments; the original file remains unchanged, and a set of new files with the same name plus an added prefix is created
 - By default, the **x** prefix is added
 - To split a file into segments, use the command **split infile**
 - To split a file into segments using a different prefix, use the command **split infile <Prefix>**



bash Script

After preparing a bash script and putting in a file, say **my_script.sh**, it can be invoked in either of two ways:

- Typing **bash my_script.sh**
- Have the first line of **my_script.sh** be **#!/bin/sh**
 - This will work with any other interpreter such as csh, perl, etc.
 - This is an exception to the rule that the character **#** is used to demark comments

Then make the script executable by doing **chmod +x my_script.sh** and then just running **./my_script.sh**

Environment Variables

- There are some special environment variables that can be used inside a script:
 - `$0` is the command name
 - `$1` `$2` ... are the command arguments
 - `$*` represents them all
 - `$@` represents them all, preserving the grouping of quoted arguments
 - `$#` gives the number of arguments
- Note that you should often enclose these variables in double quotes; i.e. you should say "`$@`" to preserve the argument grouping
- In addition, you will get syntax errors when doing comparisons if the string is empty if you do not use double quotes

Functions

- bash permits use of subprograms or functions
- Alternatively, one script can call another script, but this can make passing information a little more complex, and increases the number of files; proper use of functions can make scripts much simpler
- The basic form of a function is:

```
fun_foobar() {  
    statements  
}
```

which will not look strange to any C programmer

Functions (Cont.)

A few things to remember about functions include:

- No arguments are enclosed in the parentheses
- Functions must always be defined **before** they are used, because scripts are interpreted, not compiled
- Functions leave positional arguments unchanged, but can reset other variables; a **local var1 var2 ...** statement can make variables local

Types of Files

- Normal files (-) and directories (**d**) are the basic entries in the filesystem structure
- Symbolic links (**l**) point somewhere else on the system, such that a reference to either the link or its target is equivalent
- Named pipes (**p**), also called FIFOs, for First-In, First-Out, are used for inter-process communication
- Unix domain sockets (**s**) are used for a similar purpose, but do this through networking infrastructure, and have some more capabilities than FIFOs

Types of Files (Cont.)

- Block device nodes (**b**) and character device nodes (**c**) are used to communicate with either hardware devices, or software pseudo-devices
 - Block devices do reads and writes in predetermined fixed size chunks, and usually correspond to storage media such as disks and CD-ROMs; the I/O is generally buffered and cached
 - Character devices are addressed with streams of data, and correspond to devices such as sound cards, serial ports, etc; the I/O is generally not buffered or cached

Device nodes are almost always placed in the `/dev` directory

Everything Is a File

- In describing UNIX-like operating systems, such as Linux, one often hears the old saying: *Everything is a file*
- This fits in with the philosophy that the same basic tools can be used for a wide variety of purposes, whether you are dealing with normal files, devices, etc., and can be strung together to accomplish non-trivial tasks.
- There are exceptions to this model; for example, network devices in Linux have no corresponding filesystem entry
- Note that in Linux file extensions rarely play a defining role; the type of a file is determined by examining its contents rather than its extension

Everything Is a File

- In describing UNIX-like operating systems, such as Linux, one often hears the old saying: *Everything is a file*
- This fits in with the philosophy that the same basic tools can be used for a wide variety of purposes, whether you are dealing with normal files, devices, etc., and can be strung together to accomplish non-trivial tasks.
- There are exceptions to this model; for example, network devices in Linux have no corresponding filesystem entry
- Note that in Linux file extensions rarely play a defining role; the type of a file is determined by examining its contents rather than its extension

Access Rights

- When you do an **ls -l** as in:

```
$ ls -l a_file  
-rw-rw-r-- 1 coop aproject 1601 Mar 9 15:04 a_file
```

after the first character there are nine more which indicate the access rights granted to potential file users

- These are arranged in three groups of three:

- owner: the user who owns the file (also called user)
- group: the group of users who have access
- world: the rest of the world (also called other)

In the above listing, the user is **coop** and the group is **aproject**

read, write, execute

- Each of the triplets can have each of the following values set:
 - r: read access is allowed
 - w: write access is allowed
 - x: execute access is allowed
- If the permission is not allowed, a - appears instead of one of these characters
- Thus, in the preceding example,

```
$ ls -l a_file  
-rw-rw-r-- 1 coop aproject 1601 Mar 9 15:04 a_file
```

the user **coop** and members of the group **aproject** have read and write access, while anyone else has only read access

Activities

Applications

Terminal

S S G G O O F F X X

Thu 6 Jul, 9:18 AM

Speaker Volume

coop@x7:/tmp/TEST

File Edit View Search Terminal Help

```
x7:/tmp/TEST>umask  
0002  
x7:/tmp/TEST>rm -f afile ; touch afile ; ls -l afile  
-rw-rw-r-- 1 coop coop 0 Jul  6 09:17 afile  
x7:/tmp/TEST>umask 0022  
S x7:/tmp/TEST>rm -f afile ; touch afile ; ls -l afile  
-rw-r--r-- 1 coop coop 0 Jul  6 09:17 afile  
x7:/tmp/TEST>umask -S  
S u=rwx,g=rx,o=rx  
x7:/tmp/TEST>umask u=r,g=w,o=rw  
x7:/tmp/TEST>rm -f afile ; touch afile ; ls -l afile
```



10:56 AM

coop@x7:/tmp/TEST 2:13 / 2:28

SimpleScreenRecorder

Sound

Virtual File System (VFS)

- Linux implements a **Virtual File System (VFS)**, as do all modern operating systems
- For the most part neither the specific filesystem or actual physical media and hardware need be addressed by filesystem operations
- Furthermore, network filesystems (such as NFS) can be handled transparently
- This permits Linux to work with more filesystem varieties than any other operating system
- This democratic attribute has been a large factor in its success
- Most filesystems have full read/write access while a few have only read access and perhaps experimental write access

Virtual File System (Cont.)

- Some filesystem types, especially non-UNIX based ones, may require more manipulation in order to be represented in the VFS
- For example, variants such as vfat do not have distinct read/write/execute permissions for the owner/group/world fields; the VFS has to make an assumption about how to specify distinct permissions for the three types of user, and such behavior can be influenced by mounting operations
- Even more drastically, such filesystems store information about files in a File Allocation Table (FAT) at the beginning of the disk, rather than in the directories themselves, a basically different architectural method
- A number of newer high performance filesystems include full journaling capability

ext2, ext3, and ext4 Filesystems

- The native filesystems for Linux are **ext2** and its journaling descendants, **ext3** and **ext4**
- Each has its associated utility for formatting the filesystem (e.g. `mkfs.ext3`) and for checking the filesystem (e.g. `fsck.ext4`)
- Additionally, many parameters can be reset or tuned after filesystem creation with the program `tune2fs`; while defragmentation is generally not necessary, one can use `e4defrag` to do so
- The ext4 filesystem had its experimental designation removed with the 2.6.28 kernel release
- Pre-existing ext3 partitions could be migrated in place to ext4 without risk; with new features to be used only in subsequent file system operations
- One immediately obvious improvement is fast fsck; the speed of a filesystem check can easily go up by an order of magnitude or more

Journaling Filesystems

- Journaling filesystems recover from system crashes or ungraceful shutdowns with little or no corruption, and they do so very rapidly
- While this comes at the price of having some more operations to do, additional enhancements can more than offset the price
- In a journaling filesystem, operations are grouped into transactions
- A transaction must be completed without error, atomically; otherwise the filesystem is not changed
- A log file is maintained of transactions; when an error occurs, usually only the last transaction needs to be examined

Journaling Filesystems Freely Available Under Linux

The following journaling filesystems are freely available under Linux:

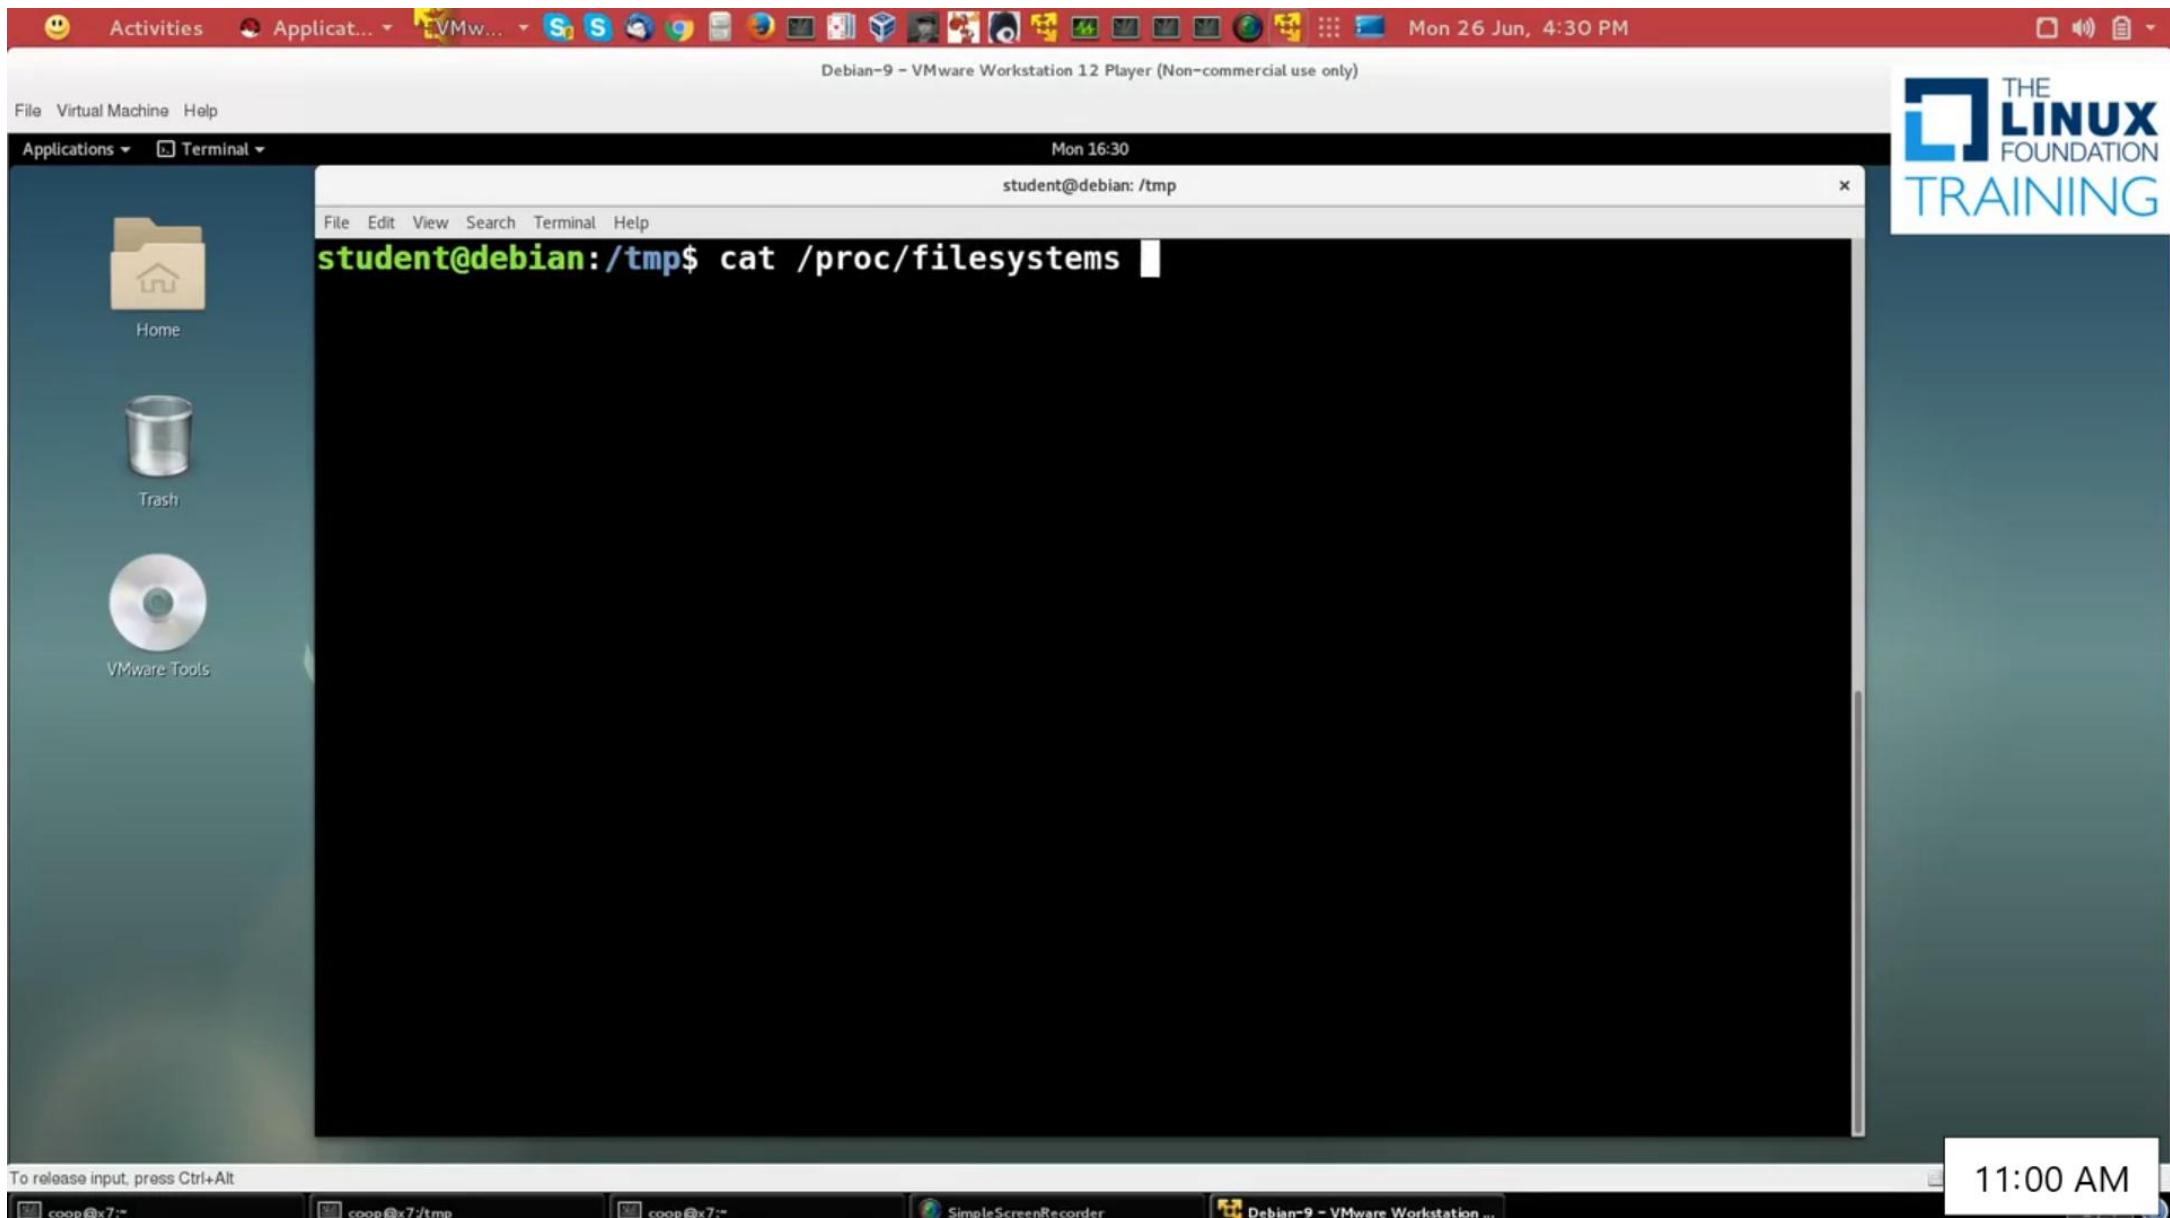
- The **ext3** filesystem is an extension of the ext2 filesystem
- The **ext4** filesystem is an extension of the ext3 filesystem, and was included in the mainline kernel first as an experimental branch, and then as a stable production feature in kernel 2.6.28 (features: extents, 48-bit block numbers, and up to 16 TB size)
- The **Reiser** filesystem was the first journaling implementation in Linux, but has lost its leadership and development has stalled
- The **JFS** filesystem is a product of IBM; has been ported from IBM's AIX OS
- The **XFS** filesystem is a product of SGI; has been ported from SGI's IRIX OS
- The **btrfs** filesystem (**B TRee** filesystem) is the most recent, is native to Linux, and has many advanced features

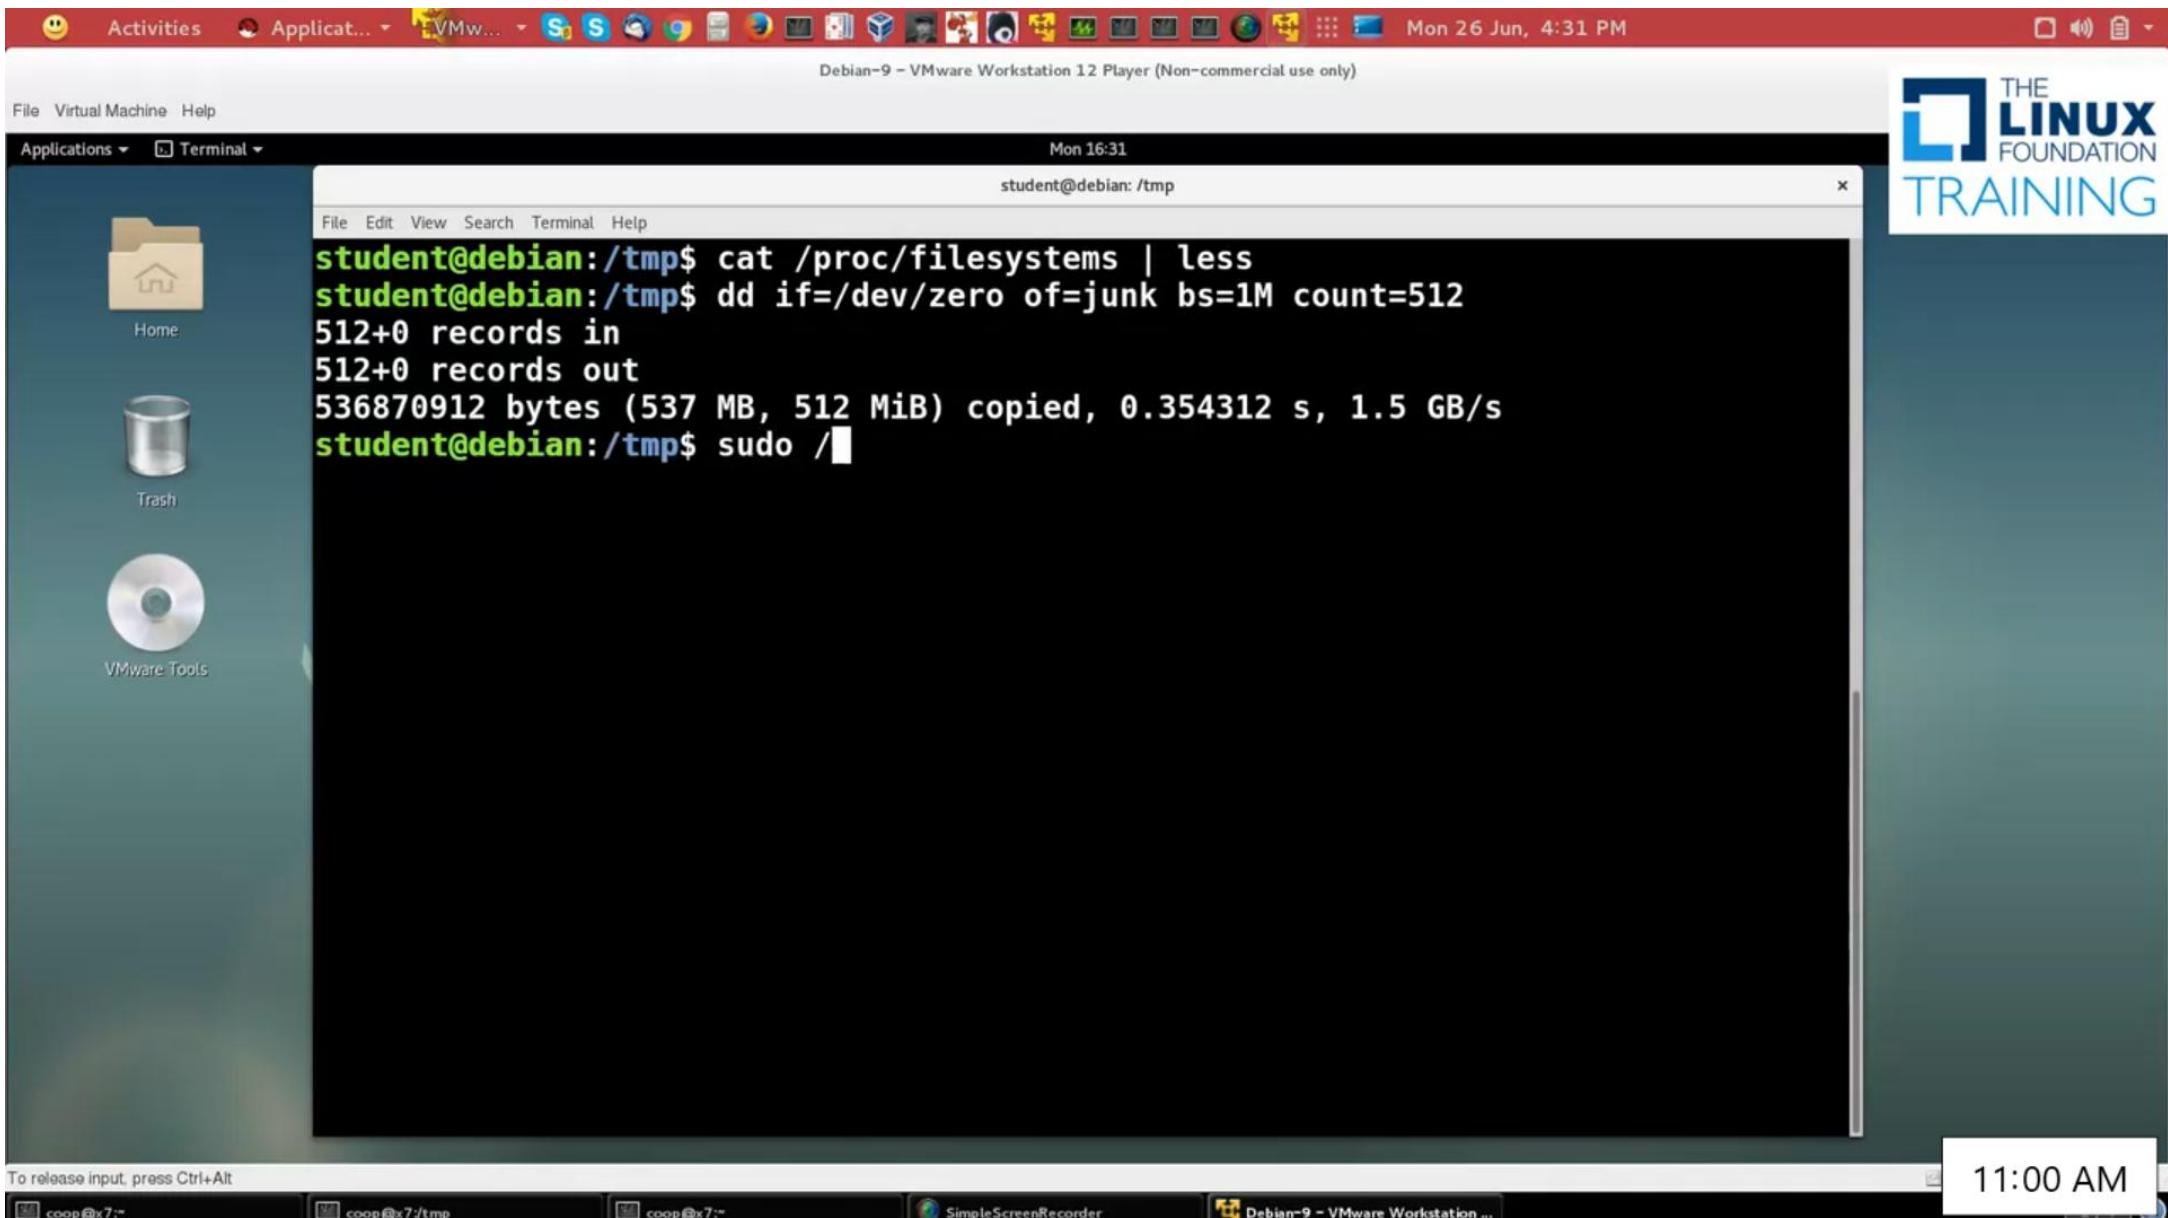
btrfs

- Both Linux developers and Linux users with high performance and high capacity or specialized needs are following the development and gradual deployment of the btrfs filesystem, which was created by Chris Mason
- While btrfs has been in the mainline kernel since 2.6.29, it has generally been viewed as experimental; although it has been used in new products
- One of the main features is the ability to take frequent snapshots of entire filesystems, or sub-volumes of entire filesystems in virtually no time
- Because btrfs makes extensive use of **COW** (Copy on Write) techniques , such a snapshot does not involve any more initial space for data blocks or any I/O activity except for some metadata updating

btrfs (Cont.)

- One can easily revert to the state described by earlier snapshots and even induce the kernel to reboot off an earlier root filesystem snapshot
- btrfs maintains its own internal framework for adding or removing new partitions and/or physical media to existing filesystems, much as LVM (Logical Volume Management) does
- Before btrfs becomes suitable for day-to-day use in critical filesystems, some tasks require finishing; for example, as of now the fsck (filesystem checking) program is read-only, which means you can find out what is wrong, but you cannot easily fix it!





THE
LINUX
FOUNDATION
TRAINING

11:00 AM

Activities Application VMw... S S Google Chrome SimpleScreenRecorder Mon 26 Jun, 4:32 PM

Debian-9 – VMware Workstation 12 Player (Non-commercial use only)

File Virtual Machine Help

Applications Terminal

student@debian: /tmp

Mon 16:32

File Edit View Search Terminal Help

meta-data=junk
 = isize=512 agcount=4, agsize=32768 blks
 = sectsz=512 attr=2, projid32bit=1
 = crc=1 finobt=1, sparse=0, rmapbt=0, refl
 = bsize=4096 blocks=131072, imaxpct=25
 = sunit=0 swidth=0 blks
 = bsize=4096 ascii-ci=0 fttype=1
 = bsize=4096 blocks=855, version=2
 = sectsz=512 sunit=0 blks, lazy-count=1
 = extsz=4096 blocks=0, rtextents=0
student@debian:/tmp\$ sudo mount junk /mnt
student@debian:/tmp\$ df -h

Filesystem	Size	Used	Avail	Use%	Mounted on
udev	1.9G	0	1.9G	0%	/dev
tmpfs	381M	12M	370M	4%	/run
/dev/sda1	18G	7.3G	9.1G	45%	/
tmpfs	1.9G	0	1.9G	0%	/dev/shm
tmpfs	5.0M	4.0K	5.0M	1%	/run/lock
tmpfs	1.9G	0	1.9G	0%	/sys/fs/cgroup
tmpfs	381M	28K	381M	1%	/run/user/117
tmpfs	381M	28K	381M	1%	/run/user/1000
/dev/sr0	56M	56M	0	100%	/media/cdrom0
/dev/loop0	509M	27M	483M	6%	/mnt

To release input, press Ctrl+Alt

coop@x7:~ coop@x7:~ coop@x7:~ SimpleScreenRecorder Debian-9 – VMware Workstation ...

11:00 AM

THE LINUX FOUNDATION TRAINING

Redundant Array of Independent Disks (RAID)

1:09 PM

- RAID (Redundant Array of Independent Disks) spreads I/O over multiple spindles, or disks; this can really increase performance in modern disk controller interfaces, such as SCSI which can perform the work in parallel efficiently
- RAID can be implemented either in software (it is a mature part of the Linux kernel) or in hardware
- If your hardware RAID is known to be of good quality it should be more efficient than using software RAID
- With the hardware implementation the operating system is actually not directly aware of using RAID; it is transparent
 - For example, three 512 GB hard drives (two for data, one for parity) configured with RAID-5 will just look like a single 1 TB disk

Features of RAID

- Three essential features of RAID are:
 - Mirroring: writing the same data to more than one disk
 - Striping: splitting of data to more than one disk
 - Parity: extra data is stored to allow problem detection and repair, yielding fault tolerance
- Thus use of RAID can improve both performance and reliability
- There are a number of RAID specifications of increasing complexity and use; the most commonly used are levels 0, 1, and 5

Logical Volume Management (LVM)

- Using **LVM** (Logical Volume Management) breaks up one virtual partition into multiple chunks, which can be on different physical volumes
- There are many advantages to using LVM; in particular it becomes really easy to change the size of the logical partitions and filesystems, to add more space, rearrange things, etc.
- One or more physical disk partitions, or **physical volumes**, are grouped together into a **volume group**
- Then, the volume group is subdivided into logical volumes which appear to the system as disk partitions and are then formatted to contain mountable filesystems

Logical Volume Management (Cont.)

- There are a variety of command line utilities tasked to create, delete, resize, etc. physical and logical volumes; fortunately, these command line utilities are not hard to use and are quite flexible
- Performance changes do occur
- There is a definite additional cost that comes from the overhead of the LVM layer; however, even on non-RAID systems, if you use striping in the setup you can achieve some parallelization improvements
- Such striping does no good if it is within the same physical disk, however



coop@c7:/tmp

--- Physical volume ---

```
PV Name          /dev/sda2
VG Name          VG2
PV Size          500.00 GiB / not usable 4.00 MiB
Allocatable      yes (but full)
PE Size          4.00 MiB
Total PE         127999
Free PE          0
Allocated PE     127999
PV UUID          sBWbb3-kUDa-oPf8-U1Qo-J8fB-CfTQ-gQ1Jr5
```

--- Physical volume ---

```
PV Name          /dev/sda5
VG Name          VG2
PV Size          500.00 GiB / not usable 4.00 MiB
Allocatable      yes
PE Size          4.00 MiB
Total PE         127999
Free PE          102399
Allocated PE     25600
PV UUID          wQlEbc-w48N-u3qz-5MAe-Q1Wm-xTf9-L3oAWf
```

c7:/tmp>



coop@c7:/tmp

File Edit View Search Terminal Help



```
VG Name          VG
VG Name          VG2
c7:/tmp>
c7:/tmp>sudo vgdisplay | grep Name
VG Name          VG
VG Name          VG2
c7:/tmp>█
```



coop@c7:/tmp



File Edit View Search Terminal Help

Act PV	3
VG Size	1.46 TiB
PE Size	4.00 MiB
Total PE	383997
Alloc PE / Size	274688 / 1.05 TiB
Free PE / Size	109309 / 426.99 GiB
VG UUID	r6lLtD-VmIN-jHCe-dEwX-4AGY-N0zZ-hXBoW8

```
c7:/tmp>sudo lvdisplay | grep Path
LV Path          /dev/VG/local
LV Path          /dev/VG/src
LV Path          /dev/VG/vms
LV Path          /dev/VG2/pictures
LV Path          /dev/VG2/dead
LV Path          /dev/VG2/iso_images
LV Path          /dev/VG2/audio
LV Path          /dev/VG2/virtual
LV Path          /dev/VG2/w7back
LV Path          /dev/VG2/P
LV Path          /dev/VG2/isabelle
LV Path          /dev/VG2/dead2
LV Path          /dev/VG2/PLAY
LV Path          /dev/VG2/newLV
```

```
c7:/tmp>sudo lvdispl■
```

SimpleScreenRecorder

coop@c7:/tmp

2 / 4

3



coop@c7:/tmp

File Edit View Search Terminal Help

c7:/tmp>**sudo lvcreate -L 4G -n newLV VG2**



coop@c7:/tmp

File Edit View Search Terminal Help

tmpfs	tmpfs	7.8G	45M	7.8G	1%	/dev/shm
tmpfs	tmpfs	7.8G	9.2M	7.8G	1%	/run
tmpfs	tmpfs	7.8G	0	7.8G	0%	/sys/fs/cgroup
/dev/sdb1	ext4	20G	11G	7.7G	58%	/
/dev/mapper/VG-local	ext4	24G	15G	7.4G	67%	/usr/local
/dev/mapper/VG-src	ext4	11G	7.5G	2.7G	74%	/usr/src
/dev/mapper/VG-vms	ext4	181G	132G	41G	77%	/VMS
/dev/loop0	squashfs	2.8G	2.8G	0	100%	/usr/src/KERNELS
/dev/mapper/VG2-virtual	ext4	345G	248G	80G	76%	/VIRTUAL
/dev/mapper/VG2-dead2	ext4	99G	60G	35G	64%	/DEAD2
/dev/mapper/VG2-dead	ext4	69G	20G	47G	30%	/DEAD
/dev/mapper/VG2-pictures	ext4	20G	13G	5.8G	70%	/PICTURES
/dev/mapper/VG2-audio	ext4	16G	9.6G	5.4G	65%	/AUDIO
/dev/mapper/VG2-iso_images	ext4	109G	44G	60G	43%	/ISO_IMAGES
tmpfs	tmpfs	1.6G	48K	1.6G	1%	/run/user/1000
/dev/mapper/VG2-newLV	ext4	3.9G	16M	3.6G	1%	/mnt

c7:/tmp>

c7:/tmp>sudo lvremove /dev/VG2/newLV

Logical volume VG2/newLV contains a filesystem in use.

c7:/tmp>sudo umount /mnt

c7:/tmp>sudo lvremove /dev/VG2/newLV

Do you really want to remove active logical volume VG2/newLV? [y/n]: yes

Logical volume "newLV" successfully removed

c7:/tmp>■

SimpleScreenRecorder coop@c7:/tmp 2 / 4 3

gcc

- `gcc` is the **GNU C** compiler; its proper name is the **GNU Compiler Collection**
- It can be invoked as `gcc` or `cc`, and can compile programs written in C, C++, and Objective C
- `g++` is the C++ compiler; it can also be invoked as `c++`
- `gcc` works closely with the GNU libc, **glibc**, and the debugger, **gdb**
- Virtually every operating system you can think of has a version of `gcc`, and it can be used for cross-compilation on different architectures

gcc (Cont.)

- gcc also forms the back end for compiler front ends in Ada95 (package `gcc-gnat`), Fortran (package `gcc-gfortran`), and Pascal (package `gcc-gpc`); i.e. first there is a translation to the C language, and then a back end (silently) invokes gcc
- The `gcc-java` package supplies `gcj` which adds support for compiling Java programs and bytecode into native code, but is no longer available on some recent Linux distributions as it is considered obsolete

Compiling Stages

- Invoking gcc actually entails a number of different programs or stages, each of which has its own **man** page, and can be independently and directly invoked; these are:

Stage	Command	Default Input	Default Output	-W Switch
Preprocessing	cpp	.c	.i	-Wp....
Compilation	gcc	.i	.s	N/A
Assembly	as	.s	.o	-Wa....
Linking	ld	.o	a.out	-Wl....

- Depending on your Linux distribution, details about the gcc installation and defaults can be found in the **/usr/lib/gcc**, **/usr/lib64/gcc** and/or **/usr/libexec/gcc** directories

LLVM

- The **LLVM** project provides an alternative C/C++ compiler called clang which is rapidly becoming a real alternative to gcc. It is already used exclusively by projects such as FreeBSD, and Mandriva Linux. Several companies are moving over to using clang to replace their current toolchains with some companies like Google and Apple moving 100% of their software development over to clang.
- At this time all IOS (iPhone/iPad) and all Android devices (phones and tablets) are now running clang compiled code.
- You can learn all about LLVM and clang at: <https://llvm.org/>.
- The ClangBuiltLinux project has continued the work of the prior LLVMLinux project upstreaming patches to the Linux kernel which enable it to be compiled with the clang compiler. Full details can be found at: <https://clangbuiltlinux.github.io/>.

Intel

- **Intel** has a mature set of compilers
- Evaluation copies can be downloaded for free, and a free non-commercial license can be obtained for learning purposes
- The Intel C compiler works well for compiling applications under Linux; it can be used to compile the kernel but it is not a trivial exercise and it is doubtful anyone is using it for this purpose in a production environment

Press Esc to exit fullscreen

Debugging with gdb

THE **LINUX** FOUNDATION

gdb

- **gdb** is the GNU debugger
- Upon launch, after processing all command line arguments and options, it loads commands from the file **.gdbinit** in the current working directory (if it exists)
- gdb allows you to step through C and C++ programs, setting breakpoints, displaying variables, etc. (actually it will work with programs in native Fortran and other languages that use gcc as a back end as well)
- In addition, gdb can properly debug multi-threaded programs
- Programs have to be compiled with the **-g** option for symbol and line number information to be available to gdb
- Note, however, you can still use gdb to get some information even if this has not been done; for instance, the **where** command often tells you exactly where the program bombed

Why Use Package Management?

- Originally, most Linux distributions were simply collections of **tarballs**, of either binaries or sources which required compilation; some distributions (**Slackware**) still work this way
- However, this method has many disadvantages:
 - Removing all files from a package can be difficult
 - One might accidentally delete a package that other packages need, or install a package that will not work because it needs other packages that have not been installed
 - A developer may lose track of what exact sources were used to build a particular binary package
 - Updates and upgrades can be very difficult, for a number of reasons:
 - Files which are no longer needed may remain on the system
 - In place upgrades done while the system is running may cause difficulties, even system crashes
 - The order of upgrading software packages may be important but not be explicitly considered
 - Software groups that require simultaneous updating may conflict with each other

RPM and APT

- Several systems have been developed in the attempt to improve on this process, these include:
 - RPM (Red Hat Package Manager): originally only Red Hat used RPM, but many other Linux distributors have used RPM, such as Fedora, and particularly SUSE and its derivatives, such as OpenSUSE
 - APT (Advanced Packaging Tool): produces deb packages; originally developed by Debian it is the foundation of Ubuntu and other Debian-derived distributions such as Linux Mint
- **alien** is a tool for converting back and forth between rpm and deb packages, or installing one kind of package on a system set up for the other; however, alien is not available on all Linux distributions

Packaging System Benefits for Developers

- Repeatable builds
- Generation of dependency data, such as what other packages are needed by a given package, and/or what other packages (or class of packages) may need a given package
- Inclusion of the pristine sources, which aids in configuration control and establishing a clear history of revision
- Full explanation of any patches that have been made to the upstream sources, together with instructions for how to proceed in building the package as well as in installing it

Packaging System Benefits for Developers

- Repeatable builds
- Generation of dependency data, such as what other packages are needed by a given package, and/or what other packages (or class of packages) may need a given package
- Inclusion of the pristine sources, which aids in configuration control and establishing a clear history of revision
- Full explanation of any patches that have been made to the upstream sources, together with instructions for how to proceed in building the package as well as in installing it

Packaging System Benefits for System Administrators and End Users

- Integrity of the installation can be verified in a uniform and rapid fashion
- Simple installation and removal methods
- Package updates and upgrades automatically preserve customized and modified configuration files
- Error-checking on install and remove (erase) ensures needed resources are not obliterated
- Users and administrators can do queries on matters including identifying what files are part of a package, or the inverse process of asking what package (if any) a given file is part of

Maintenance

- Maintenance of software by using packaging systems is an essential task of any Linux distribution
- Every other function that must be accomplished requires coherent packaging, including:
 - Keeping a healthy bi-directional connection to the upstream developers, both incorporating patches from upstream as well as sending patches to the upstream maintainers
 - Establishing and deploying clear policies for dealing with configuration files, etc., as packages are updated
 - Updating and upgrading packages in timely fashion, both to fix outstanding bugs and security holes, and to incorporate new features
 - Most importantly, perhaps, ensuring proper package dependencies, including control the order of installation and removal, requiring various packages to be dealt with as one group, etc.
 - Maintaining (multiple) repositories, ensuring that they are secure and authoritative, as well as complete

Maintenance (Cont.)

- Linux distributions include software from many sources
- Maintaining coherence, ensuring proper licensing compliance, etc., is actually a very difficult task.
- The package maintainer may be an employee of the distribution (especially for commercial distributions such as Red Hat or Ubuntu) or a “volunteer” for other distributions such as Debian or Fedora
- In either case, the maintainer has to think long-term, how to really do a good job on the package building so that updates, upgrades, use of different systems, can be done rather cleanly with time

Maintenance (Cont.)

- Linux distributions include software from many sources
- Maintaining coherence, ensuring proper licensing compliance, etc., is actually a very difficult task.
- The package maintainer may be an employee of the distribution (especially for commercial distributions such as Red Hat or Ubuntu) or a “volunteer” for other distributions such as Debian or Fedora
- In either case, the maintainer has to think long-term, how to really do a good job on the package building so that updates, upgrades, use of different systems, can be done rather cleanly with time

